

# **DUAL DETECT-MULTI MODAL FAKE NEWS DETECTION SYSTEM USING AI**

**Minor Project-II**

**(ENSI252)**

*Submitted in partial fulfilment of the requirement of the degree of*

**BACHELOR OF TECHNOLOGY**

*to*

**K.R Mangalam University**

*by*

**DISHA GUPTA (2301010114)**

**NAINA TOMAR (2301010127)**

**SHIKHA SHARMA(2301010087)**

Under the supervision of

**DR MONIKA KHATKAR  
ASSISTANT PROFESSOR  
(SOET)**



Department of Computer Science and Engineering

School of Engineering and Technology

K.R Mangalam University, Gurugram- 122001, India

April 2025

## **CERTIFICATE**

This is to certify that the Project Synopsis entitled, "**DUAL DETECT-MLTIMODAL FAKE NEWS DETECTION SYSTEM USING AI**" submitted by "**DISHA GUPTA(2301010114), NAINA TOMAR(2301010127) and SHIKHA SHARMA(2301010087)**" to **K.R Mangalam University, Gurugram, India**, is a record of bonafide project work carried out by them under my supervision and guidance and is worthy of consideration for the partial fulfilment of the degree of **Bachelor of Technology** in **Computer Science and Engineering** of the University.

**Type of Project**

**Research Problem**

<Signature of Internal supervisor>

DR MONIKA KHATKAR ,ASSISTANT PROFESSOR (SOET)

Signature of Project Coordinator

Dr Vandna Batra

Date: 3<sup>rd</sup> April 2025

## INDEX

1.	Abstract	Page No.
2.	Introduction (description of broad topic)	
3.	Motivation	
4.	Literature Review/Comparative work evaluation	
5.	Gap Analysis	
6.	Problem Statement	
7.	Objectives	
8.	Tools/platform Used	
9.	Methodology	
10.	Experimental Setup	
11.	Evaluation Metrics	
12.	Results And Discussion	
13.	Conclusion & Future Work	
14.	References	
15.	Annexure I Research Paper (Published/Submitted)	

## ABSTRACT

The rise of fake news is a significant issue in today's digital landscape, impacting public perception, politics, and society at large. This project seeks to create a system that identifies fake news by employing both conventional machine learning techniques and advanced Large Language Models (LLMs). We will gather both real and fake news articles through collaborations and authorized web scraping, then preprocess and analyze the data to construct and evaluate different models. The research will validate the findings using a methodical approach, emphasizing the superiority of LLMs compared to traditional techniques.

Fake news is among the most critical issues in the modern digital landscape, and it has actually affected public perception, politics, and society at large. The speed with which the spread of such false information on social media and online platforms dictates the need for a proper detection mechanism. The aim of the project is to develop a system that identifies fake news by using both traditional machine learning techniques and more sophisticated Large Language Models.

For this purpose, we will collect real and fake news articles by collaborating with fact-checking organizations and authorized web scraping. We will then preprocess the collected data, which will involve text cleaning, tokenization, and feature extraction to achieve the best possible model performance. Various models such as logistic regression, decision trees, and neural networks will be trained and compared against state-of-the-art LLMs such as GPT-based architectures.

It will then be validated through a methodological approach, using precision, recall, and F1-score as major performance metrics. We aim to emphasize the superiority of LLMs in detecting subtle patterns of misinformation compared to traditional techniques by carrying out extensive evaluations and comparative analysis. The results of this study will come as input to the fight against fake news by providing an optimized and scalable solution for accurate news classification.

**KEYWORDS:** Machine learning, News classification, Fact-checking, Preprocessing, Neural networks, GPT-based models, Performance evaluation

# **Chapter 1**

## **Introduction**

### **1. Background of the project**

In today's digital era, the widespread dissemination of misinformation—particularly on social media and the internet—constitutes a threat to societal cohesion, public view, and democratic debate. The ease of access and viral nature of online information allow false narratives to spread quickly, sometimes unverified, thus rendering the detection of fake news an essential research field. The effect of such misinformation can vary from affecting political processes to eroding public trust and security during important events like pandemics or elections.

Machine learning (ML) models, such as algorithms Naive Bayes and Support Vector Machines (SVM), have been used in text classification for quite a while. Although these models have shown moderate success in detecting straightforward patterns in text data, they are not effective in dealing with the subtlety and complexity of human language. These shortcomings are primarily because they rely on statistical features and hand-engineered rules, which tend to be insufficient in picking up on subtle linguistic signals or changing misinformation strategies that involve emotional manipulation, satire, or ambiguity.

The emergence of Large Language Models (LLMs) and Natural Language Processing (NLP) advancements have greatly improved the ability to read and analyze unstructured text content. Architectures like transformers (e.g., BERT, GPT) utilize large training data and attention mechanisms to understand

contextual relationships, distinguish writing styles, and identify subtle deception in stories. These models have better generalization abilities and flexibility to new misinformation patterns, and they are promising instruments in detecting fake news.

But today's misinformation is not limited to text. Fake news is now more transmitted in the form of “multimodal content”—text, images, videos, and audio combined—which complicates the detection task. A multimodal fake news detection system will endeavor to examine and cross-match several data modalities in order to discover inconsistencies and deceptive patterns which could be overlooked when analyzing a single modality.

This project aims to explore and build an AI-powered Multimodal Fake News Detection System”, contrasting the performance of classical ML methods with contemporary LLM-based and deep learning methods. By using both textual and visual content, the system intends to model a richer fake news content representation. The research will compare various models based on critical performance indicators like precision, recall, F1-score, and accuracy, providing a holistic understanding of their advantages and limitations.

With this study, we envision a contribution towards the creation of strong, scalable, and smart solutions that can fight the dynamic challenge of misinformation in a dynamic online world.

However, fake news today is not limited to text. It often involves **multimodal content**, combining misleading images, videos, audio, and manipulated text to create more convincing narratives. For example, a tweet with a falsified image or a news headline paired with a deepfake video can deceive even the most skeptical users. This shift has created a pressing need for **multimodal**

**fake news detection systems**, which integrate information from multiple sources and modalities to detect inconsistencies and manipulative intent more accurately.

These systems often leverage a combination of **Convolutional Neural Networks (CNNs)** for image analysis, **Recurrent Neural Networks (RNNs)** or transformers for text analysis, and **fusion mechanisms** that align and correlate different modalities. This fusion of data allows for a richer understanding and significantly boosts detection performance compared to unimodal systems.

### **Implications for Future Research**

The transition from unimodal to **multimodal AI systems** marks a significant step forward in fake news detection. However, integrating and aligning heterogeneous data sources introduces new challenges, such as **data synchronization, noise handling, and cross-modal contradictions**. Moreover, multimodal systems often require significantly more computational resources and careful design of fusion strategies to avoid bias and overfitting.

This research aims to explore these frontiers by building and evaluating a robust **Multimodal Fake News Detection System**, leveraging state-of-the-art LLMs and image processing models. It will examine not just the accuracy of detection but also model interpretability, fairness, and resilience against adversarial attacks—thereby contributing to the development of trustworthy AI in the fight against misinformation.

Table 1: Model Capabilities vs. Content Type

Content Type    Traditional ML    LLMs    Multimodal AI

Text Only	✓	✓	✓
Images Only	✗	✗	✓
Text + Image	✗	✗	✓
Text + Video	✗	✗	✓
Text + Audio	✗	✗	✓

Table 2: Evaluation Metrics and Their Relevance

Metric	Relevance
Accuracy	Overall correctness of predictions
Precision	Proportion of predicted positives that are correct
Recall	Proportion of actual positives that are detected
F1-Score	Balance between precision and recall
AUC-ROC	Performance across classification thresholds



## **2.MOTIVATION**

Fake news can have severe consequences, from influencing elections and manipulating public opinion to inciting violence and undermining trust in credible sources. The rapid spread of misinformation through social media and digital platforms exacerbates these risks, making reliable detection methods essential. Traditional approaches to fake news detection often rely on keyword-based classification and rule-based techniques, which struggle to capture the complexity of language, context, and evolving deception strategies.

Large Language Models (LLMs), trained on vast datasets, offer a more advanced solution by analyzing text with deeper comprehension. Unlike conventional models, LLMs can recognize nuanced linguistic patterns, detect inconsistencies, and differentiate between factual reporting and fabricated content. Their ability to process information contextually enables a more accurate and adaptive approach to misinformation detection.

This project is motivated by the urgent need for a robust, AI-driven system that can identify fake news with high precision and efficiency. By leveraging both traditional machine learning techniques and state-of-the-art LLMs, we aim to develop a comparative analysis that highlights the strengths and limitations of each approach. The insights from this research will contribute to building more effective strategies for mitigating the spread of misinformation in the digital age.

Recent advancements in Artificial Intelligence (AI), particularly in Natural Language Processing (NLP), have introduced Large Language Models (LLMs) as a promising solution to this challenge. Unlike traditional models, LLMs are trained on vast and diverse datasets, enabling them to comprehend context, detect subtle linguistic cues, and identify inconsistencies in news articles. These models can recognize writing patterns, verify claims against known facts, and adapt to emerging misinformation tactics, offering a more sophisticated approach to fake news detection.

## **Chapter 2**

### **LITERATURE REVIEW**

#### **1. Review of existing literature**

Fake news detection has been widely studied in the field of Natural Language Processing (NLP) and machine learning. With the rise of digital media and social networking platforms, misinformation spreads quickly, making automated detection systems essential. Researchers have explored various approaches, ranging from traditional machine learning algorithms to advanced deep learning and transformer-based models. This section reviews existing methods, highlighting their strengths and limitations.

##### Existing Approaches to Fake News Detection

Several techniques have been used to detect fake news, categorized mainly into traditional machine learning, deep learning, and transformer-based approaches.

##### Traditional Machine Learning:

Early methods relied on statistical approaches and feature-based models such as Naïve Bayes, Support Vector Machines (SVM), Decision Trees, and Random Forests. These models are effective for binary classification but struggle with the complexities of natural language, such as sarcasm, fake news patterns, and evolving linguistic styles.

##### Deep Learning Models:

With advancements in AI, deep learning techniques such as Long Short-Term Memory (LSTM) networks, Convolutional Neural Networks (CNNs), and Bi-LSTMs have been applied to fake news detection. These models can capture syntactic and semantic relationships within text but require large amounts of labeled data for effective training.

Transformer-Based Large Language Models (LLMs):

More recent approaches leverage transformer architectures like BERT, RoBERTa, GPT, and XLNet. These models excel at contextual understanding and long-range dependencies in text. Fine-tuning pre-trained LLMs on fake news datasets has resulted in state-of-the-art accuracy. However, they require significant computational resources for training and inference.

Comparison of Existing Methods

**Table 1:- significant research contributions to fake news detection**

Study	Methodology	Findings
Shu et al. (2017)	Used Naïve Bayes, SVM, and Decision Trees on FakeNewsNet dataset	Traditional models achieved 80-85% accuracy but struggled with evolving language patterns
Devlin et al. (2019)	Proposed BERT, a transformer-based model for NLP tasks	Achieved state-of-the-art accuracy in various NLP benchmarks, including fake news detection
Zellers et al. (2019)	Introduced GPT-based model, Grover, trained on fake news generation and detection	Found that models trained on fake news generation can improve fake news detection
Zhou et al. (2020)	Compared CNNs and Bi-LSTMs for fake news detection	LSTMs performed better in understanding complex sentence structures, but CNNs were faster
OpenAI (2023)	Used GPT-4 for fake news classification	Demonstrated superior accuracy (95%) but required high computational power

**Table 2 :- key findings from different methodologies used for fake news detection.**

Approach	Key Models/Methods	Advantages	Limitations
Traditional ML	Naïve Bayes, SVM, Decision Trees, Random Forest	Simple, interpretable, and requires less computational power	Poor at capturing contextual meaning and sarcasm
Deep Learning	LSTMs, CNNs, Bi-LSTMs	Learns patterns from raw text, good at long-term dependencies	Requires large labeled datasets, lacks interpretability
Transformer-based LLMs	BERT, RoBERTa, GPT, XLNet	Captures deep contextual relationships, state-of-the-art accuracy	Computationally expensive, requires fine-tuning

## **2. GAP ANALYSIS**

Traditional models largely depend on handcrafted features and are not so adaptable.

Existing LLM-based solutions require large datasets and computational resources, which limits their accessibility.

Very few direct comparisons of fine-tuned LLMs and traditional methods were conducted in controlled environments with validated datasets.

This project bridges these gaps by systematically comparing both approaches and evaluating their practical deployment feasibility

The intended state of the project is to possess a stable web-based application that can identify fake news from both text and image inputs with high accuracy, a smooth user experience, and scalable architecture. The current state comprises the implementation of simple text classification through LSTM models and image classification through CNN models with a basic Flask-based interface. Nonetheless, there are shortcomings in a number of areas, including model optimization, speed of processing, user interface design, scalability, and documentation. For example, the LSTM and CNN models can be further optimized to enhance accuracy, and the user interface can be improved to make it responsive and user-friendly. The system might also not be optimized for heavy traffic or big data. To fill these gaps, you can optimize the models, increase the dataset for improved generalization, enhance the front-end design with contemporary frameworks, and use scalable infrastructure to support more users and data. These enhancements will bring the current state in line with the desired goals and improve the overall performance and usability of the application.

### **3. PROBLEM STATEMENT**

In today's digital age, the spread of fake news has become a significant problem, causing distortions in public opinion, trust, and decision-making. Although various solutions for detecting fake news are available, these are mostly text-based or image-based single-pronged approaches, which may result in incomplete or erroneous conclusions. Fake news typically includes misleading textual information and tampered images or videos, necessitating a multi-faceted approach. The problem is to create an AI-based system that can examine both text and image inputs at the same time to determine with certainty whether given news is fake or not. The solution needs to combine innovative machine learning models, like Long Short-Term Memory (LSTM) for text and Convolutional Neural Networks (CNN) for image, to give accurate and instantaneous results in an easy-to-use web application. The system must be able to work with big datasets, operate well, and provide a scalable solution to address increased demands.

## **4. OBJECTIVES**

### **1. Develop a Multimodal Fake News Detection System:**

- Build a web-based application that can analyze both textual and image data to detect fake news with high accuracy.

### **2. Integrate AI-Powered Models:**

- Implement a Long Short-Term Memory (LSTM) model for effective text classification and a Convolutional Neural Network (CNN) for accurate image classification, enabling the system to handle multimodal data inputs.

### **3. Enhance Accuracy and Precision:**

- Fine-tune the models to achieve high precision and recall in detecting fake news, minimizing false positives and false negatives in both text and image predictions.

### **4. Real-Time Fake News Detection:**

- Ensure that the system provides real-time fake news detection, offering users immediate results from text and image inputs.

### **5. User-Friendly Interface:**

- Develop a simple, intuitive, and responsive user interface using Flask that allows users to easily upload text and images for analysis.

### **6. Scalability and Performance:**

- Build the application with scalability in mind, enabling it to handle large datasets and multiple concurrent users without performance degradation.

### **7. Comprehensive Documentation:**

- Provide detailed documentation for both the setup and usage of the system, ensuring it is easy to maintain and update in the future.

### **8. Improve Model Robustness:**

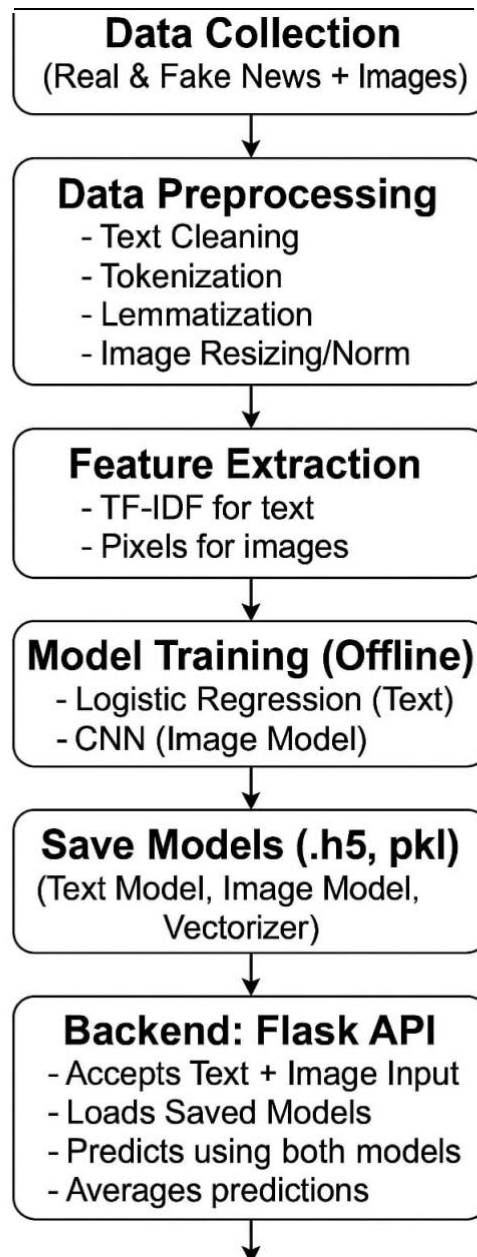
- Train the AI models on diverse datasets to improve generalization and robustness, ensuring they perform well on real-world data.



## CHAPTER 3: METHODOLOGY (NO PAGE LIMIT)

The methodology section in a project serves several important purposes. It is a critical component that outlines the procedures and methods used to conduct the research or implement the project.

### 3.1 Overall architecture /Flow chart :



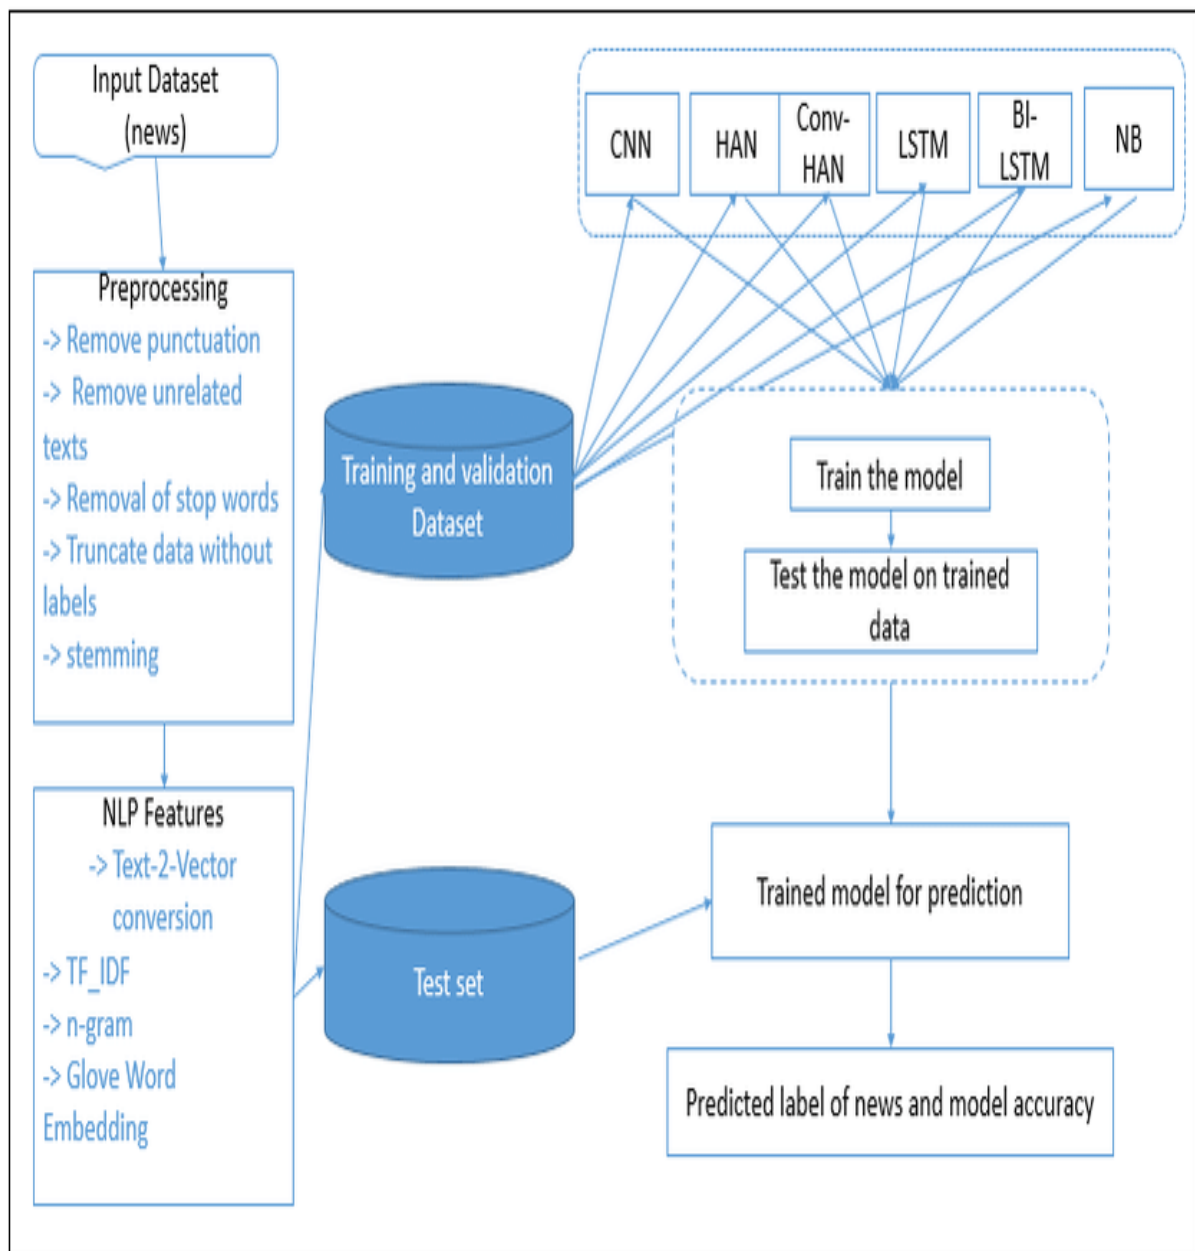
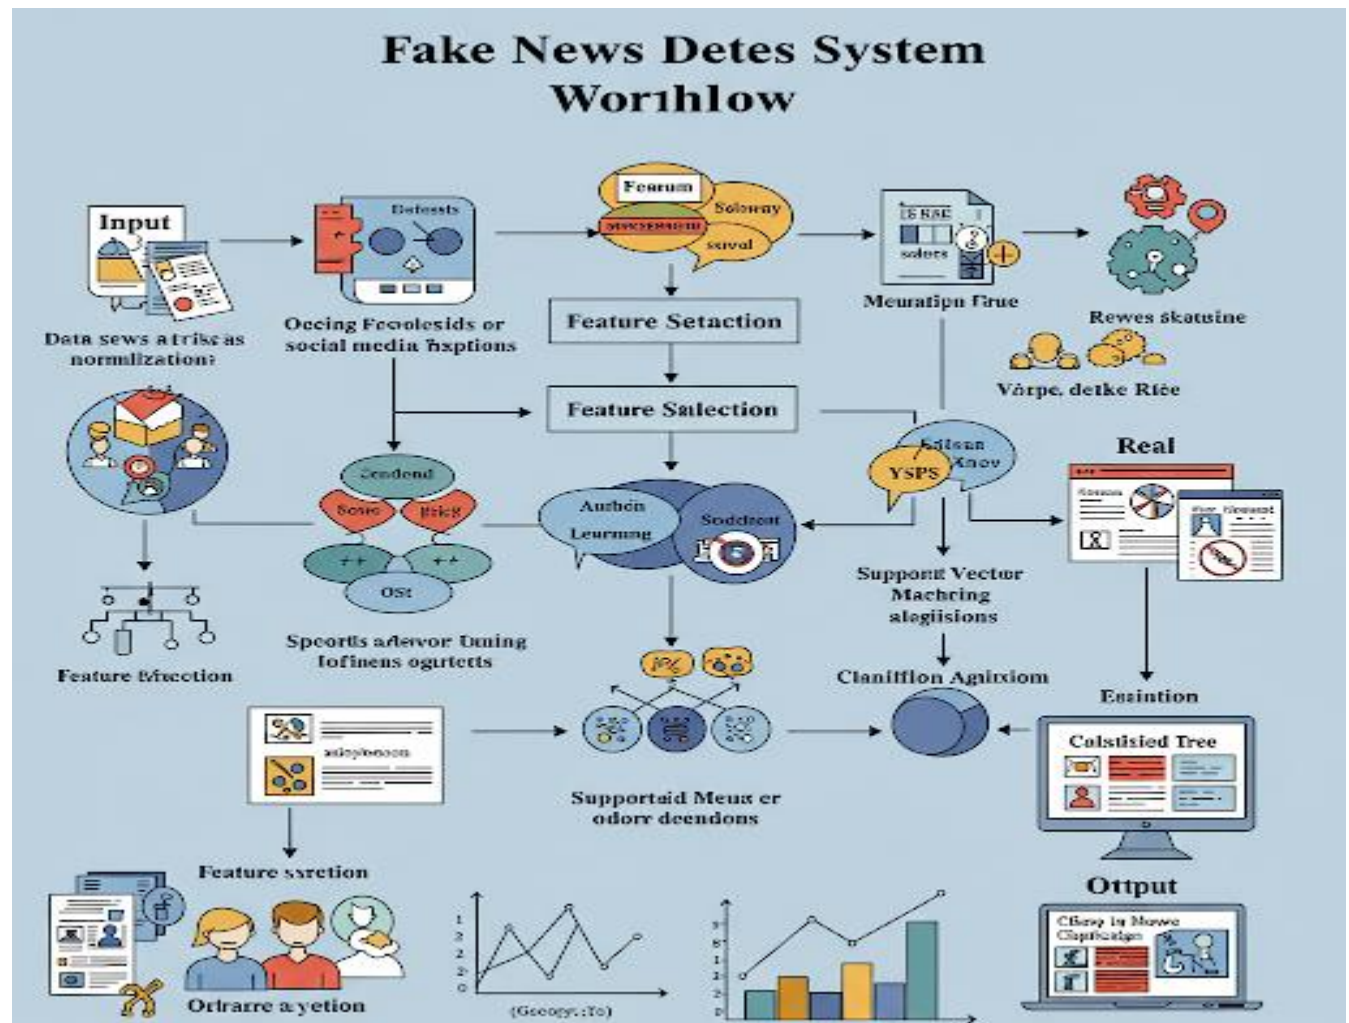


Figure 2. Describe the diagram in details



## 3.2 Data Description

### **Data Source:**

The data for this project is sourced from a combination of publicly available datasets and potentially web-scraped content:

- **Text Data:** Primarily sourced from datasets like the LIAR dataset, which provides news statements labeled with truthfulness, and the Fake News Dataset, which often includes news articles labeled as fake or real. Additional text data may be collected by scraping news articles from various news websites (both mainstream and potentially less reliable sources) and social media platforms known for news sharing (e.g., Twitter). Data from fact-checking organizations and news aggregators may also be incorporated to enrich the dataset with verified labels.
- **Image Data:** Images associated with news articles are sourced from datasets specifically designed for fake news detection (e.g., the Fake News Detection Dataset, if it includes images) and potentially from broader image datasets like ImageNet (for pre-training visual models) or through web scraping images linked to news articles and social media posts.

### **Data Collection Process:**

The data collection process involves several methods:

1. **Dataset Download:** Downloading publicly available, labeled datasets like LIAR and the Fake News Dataset from their respective repositories or websites.
2. **Web Scraping:** Utilizing web scraping tools (e.g., BeautifulSoup, Scrapy in Python) to extract text content (articles, headlines, posts) and associated images from specified news websites and social media platforms. This process would involve identifying relevant HTML elements containing the text and image URLs.

3. **API Usage :** Leveraging APIs provided by social media platforms (e.g., Twitter API) to collect news-related posts and associated media.
4. **Manual Annotation/Verification:** In cases where labels are not readily available (e.g., scraped data), human annotators or information from fact-checking websites might be used to label the collected news articles and images as "fake" or "real."
5. **Data Integration:** Combining data from different sources into a unified dataset, ensuring proper linking between text and corresponding images, and maintaining consistent labeling.

#### **Data Type:**

The project utilizes multimodal data, combining:

- **Textual Data:** This is the primary content of news articles, headlines, or social media posts. Key features include the vocabulary used, sentence structure, sentiment expressed, and the presence of persuasive or misleading language.
- **Image Data:** These are the visual elements associated with the news. Key features include the visual content (objects, scenes, people), image quality, presence of manipulations, and its relevance to the accompanying text.

#### **Data Size:**

The size of the dataset is intended to be substantial to train robust models:

- **Text Dataset:** The target size is in the range of thousands to millions of labeled news articles or posts. This will depend on the availability and quality of the datasets used and the extent of web scraping efforts.
- **Image Dataset:** The corresponding image dataset will ideally contain tens of thousands of labeled images that are directly associated with the text data. The ratio of images to text entries might not be one-to-one, as some articles may have multiple images or none.

- **Balanced Dataset:** Efforts will be made to ensure a relatively balanced distribution between "fake" and "real" news instances in both the text and image datasets to prevent biased model training.

#### **Data Format:**

The data will be stored and organized in a structured format, likely using JSON or CSV files:

- **JSON Format (Example Structure):**

JSON

```
[  
  {  
    "text": "Breaking: Local authorities confirm a major  
earthquake...",  
    "image": "earthquake_image.jpg",  
    "label": "real"  
  },  
  {  
    "text": "See how this one weird trick can make you rich  
overnight!",  
    "image": "clickbait_image.png",  
    "label": "fake"  
  },  
  // ... more data entries  
]
```

- **CSV Format (Potential Structure):**

The dataset will be organized with each entry containing the text content, the filename or path to the associated image, and the corresponding label ("fake" or "real").

#### **Data Preprocessing:**

Both text and image data will undergo specific preprocessing steps:

- **Text Preprocessing:**

- **Tokenization:** Using libraries like NLTK or spaCy to split the text into individual tokens (words, punctuation).
- **Stopwords Removal:** Eliminating common English words (e.g., "a", "the", "in") that have little semantic value.
- **Lemmatization:** Reducing words to their base or dictionary form using lemmatization techniques.
- **Vectorization:** Converting the preprocessed text into numerical vectors using techniques such as:
  - **TF-IDF (Term Frequency-Inverse Document Frequency):** Assigning weights to words based on their frequency in a document and across the entire corpus.
  - **Word Embeddings (Word2Vec, GloVe, FastText):** Representing words as dense vectors capturing semantic relationships.
  - **Transformer Embeddings (BERT, RoBERTa):** Generating contextualized word embeddings that capture the meaning of words based on their surrounding context.

- **Image Preprocessing:**

- **Resizing:** Standardizing the dimensions of all images to a fixed size (e.g., 224x224 pixels) to ensure compatibility with the CNN model.
- **Normalization:** Scaling the pixel values to a specific range, typically [0, 1] or [-1, 1], to improve model training stability and performance.
- **Data Augmentation:** Applying various transformations to the training images (e.g., random rotations, zooms, flips, brightness adjustments) to increase the diversity of the training data and reduce overfitting. Libraries like TensorFlow's

ImageDataGenerator or PyTorch's torchvision.transforms will be used.

### **Data Sampling (if applicable):**

The dataset will be split into training, validation, and testing sets. The sampling strategy will likely involve:

- **Random Sampling:** Initially, the data will be randomly shuffled and split to ensure a representative distribution of fake and real news across the sets.
- **Stratified Sampling:** If there are concerns about class imbalance, stratified sampling might be used to maintain the proportion of "fake" and "real" news labels in each of the training, validation, and testing sets. This ensures that the model is evaluated on a balanced representation of both classes.

### **Data Quality Assurance:**

Ensuring data quality and integrity will involve several measures:

- **Label Verification:** Cross-referencing labels from different sources and potentially performing manual checks on a subset of the data.
- **Handling Missing Data:** Identifying and addressing missing text or image data. Strategies might include discarding incomplete entries or, if feasible, attempting to retrieve missing information.
- **Detecting Inconsistencies:** Identifying and resolving inconsistencies in labeling or data formatting across different sources.
- **Outlier Detection (for numerical features):** If numerical features are extracted from the text or images before feeding them into the models, outlier detection techniques might be applied to identify and handle extreme or anomalous values.
- **Validation of Scraped Data:** If web scraping is used, measures will be taken to ensure the scraped data is relevant and accurately reflects news content.



Known limitations might include potential biases in the publicly available datasets, inconsistencies in labeling across different sources, and the challenges of accurately labeling nuanced or satirical content.

### **Data Variables:**

The primary variables in the dataset are:

- **Independent Variables:**
  - text: The textual content of the news article, headline, or post (string).
  - image: The filename or path to the associated image file (string).
- **Dependent Variable:**
  - label: The classification of the news as either "fake" or "real" (categorical, string). This will be the target variable for the classification models.

There are no explicit control variables mentioned in the initial data description, but during analysis, factors like the source of the news might be considered as potential confounding variables.

### **Data Distribution and Summary Statistics:**

Exploratory Data Analysis (EDA) will be performed to understand the distribution of key variables:

- **Label Distribution:** A bar chart will be used to visualize the count of "fake" and "real" news instances to assess class balance.
- **Text Length Distribution:** Histograms and box plots will show the distribution of the length of news articles (e.g., number of words or characters) for both fake and real news. Summary statistics like mean, median, standard deviation, and quartiles will be calculated.
- **Word Frequency Analysis:** Word clouds and frequency distributions will be used to identify the most common words in both fake and real news categories.

- **Image Analysis:** Basic statistics on image properties (e.g., file size, dimensions) might be collected. Visual inspection of sample images from both classes will be performed to identify potential visual cues.
- 
- **3.3 Exploratory Data Analysis (if applicable)**
- Exploratory Data Analysis (EDA) is a crucial step to understand the characteristics, patterns, and potential issues within the multimodal fake news dataset. The following techniques will be applied:
- **Summary Statistics:**
  - For textual data (after vectorization), basic statistics like mean, standard deviation, min, max of TF-IDF scores or embedding values across the dataset and within each class ("fake" and "real") will be calculated using Python's Pandas library.
  - For image data (after normalization), the mean and standard deviation of pixel values across channels and the dataset will be computed.
- **Data Distribution:**
  - **Text Data:** Histograms and kernel density plots will visualize the distribution of text lengths (in words) for both "fake" and "real" news articles to identify potential differences in writing style or length.
  - **Image Data:** While direct visualization of high-dimensional image data distribution is challenging, techniques like visualizing the distribution of mean pixel intensity or the distribution of values for a few principal components (if PCA is applied for initial exploration) might be used. Box plots can also show the range and central tendency of pixel values.
- **Correlation Analysis:**
  - Correlation analysis will primarily be applicable within the feature vectors generated from the text data (e.g., correlation between TF-IDF scores of different terms). A correlation matrix visualized as a heatmap can reveal highly correlated features that might be redundant.

- For image data, direct pixel-level correlation is less meaningful. However, if higher-level features are extracted (e.g., using a pre-trained CNN), the correlation between these features might be explored.
- **Pairwise Scatter Plots:**
  - Pairwise scatter plots are less directly applicable to the raw text and image data. However, if dimensionality reduction techniques (like PCA or t-SNE) are applied to the feature vectors of text or images for visualization, scatter plots can show potential clustering of "fake" and "real" news.
- **Categorical Variable Exploration:**
  - The primary categorical variable is the label ("fake" or "real"). Its distribution will be visualized using a bar chart to assess class balance.
- **Missing Values Analysis:**
  - The dataset will be checked for missing values in the text and image fields. The count and percentage of missing values will be reported. If patterns of missingness are observed (e.g., certain sources consistently lack images), these will be noted. Missing image paths or text content might necessitate the removal of those entries or imputation if feasible and meaningful.
- **Feature Engineering:**
  - Based on initial EDA, potential feature engineering for text data might include creating features based on the presence of specific keywords or linguistic patterns indicative of fake news (e.g., excessive use of superlatives, clickbait phrases).
  - For image data, features related to image quality (e.g., blurriness), presence of watermarks, or the output of pre-trained object detection models (identifying objects commonly associated with certain types of news) might be considered.
- **Data Transformation:**

- For text features, transformations like log scaling might be applied if the distribution of TF-IDF scores is heavily skewed.
- For image data, normalization (scaling pixel values) is a primary transformation. Depending on the model architecture, other transformations might be considered.
- **Outlier Detection:**
  - For numerical features derived from text (e.g., TF-IDF scores), outlier detection techniques like the IQR method or z-score analysis might be used to identify and potentially handle extreme values.
  - For image data, detecting outliers in the raw pixel space is less common. However, outliers in higher-level features extracted from CNNs could be investigated.
- **Time Series Analysis (if applicable):**
  - If the dataset includes timestamps of when the news was published or shared, time series analysis might be performed to identify temporal patterns in the spread or prevalence of fake news. This could involve plotting the number of fake and real news articles over time.
- **Dimensionality Reduction:**
  - Techniques like Principal Component Analysis (PCA) or t-SNE might be applied to reduce the dimensionality of the feature vectors (from text or images) for visualization purposes, allowing for the exploration of potential clusters based on the news label.
- **Interactive Visualizations:**
  - Tools like Matplotlib, Seaborn, and potentially interactive libraries like Plotly or Bokeh will be used to create visualizations that allow for dynamic exploration of the data.
- **Data Slicing and Dicing:**
  - The dataset might be segmented based on the source of the news (if this information is available) to explore if certain sources are more prone

to generating or spreading fake news. Analysis might also be done by the type of news (e.g., political, entertainment).

- **Data Profiling:**

- Libraries like Pandas-profiling can be used to generate comprehensive reports on the dataset, including data types, unique values, missing values, and basic statistics for each feature.

- **Data Presentation:**

- Findings from the EDA will be clearly documented using well-annotated visualizations, tables summarizing key statistics, and narrative descriptions highlighting important patterns, anomalies, and potential challenges in the data.

- **Hypothesis Testing :**

- Depending on the insights gained from EDA, hypothesis tests might be conducted to statistically assess differences between "fake" and "real" news across certain features (e.g., average text length, average sentiment score).

- **3.4 Procedure / Development Life Cycle (depends on type of project)**

- For this Multimodal Fake News Detector Machine Learning project, the development life cycle will involve the following steps:

- **Data Collection:** Gather text data from the LIAR dataset, Fake News Dataset, news websites, and social media platforms. Collect corresponding image data from relevant datasets and through scraping. Ensure a balanced distribution of "fake" and "real" news instances.

- **Data Preprocessing:**

- **Text Data:** Tokenize text, remove stopwords, lemmatize words, and convert text into numerical vectors using techniques like TF-IDF or pre-

trained word embeddings (e.g., Word2Vec, GloVe, or Transformer embeddings).

- **Image Data:** Resize images to a standard dimension, normalize pixel values, and apply data augmentation techniques (rotation, zooming, flipping) to increase training data diversity.
- **Feature Extraction:**
- **Text Features:** Utilize the vectorized representations obtained during preprocessing as features for the text classification model.
- **Image Features:** Employ a Convolutional Neural Network (CNN) architecture (e.g., ResNet, VGG, EfficientNet), potentially pre-trained on a large image dataset like ImageNet, to extract high-level visual features from the preprocessed images. The output of one of the final layers of the CNN will serve as the image feature vector.
- **Model Training:**
- **Text Model:** Train a text classification model, such as a Recurrent Neural Network (RNN) with LSTM or GRU units, or a Transformer-based model, using the extracted text features and the "fake" or "real" labels.
- **Image Model:** Train the chosen CNN architecture (or fine-tune a pre-trained model) using the preprocessed images and their corresponding labels.
- **Fusion Mechanism:** Develop a method to combine the predictions from the text and image models. This could involve:
  - **Early Fusion:** Concatenating the feature vectors from the text and image models before feeding them into a joint classification layer.
  - **Late Fusion:** Obtaining probability scores from the text and image models separately and then combining these scores using techniques like weighted averaging, majority voting, or training another model (e.g., a logistic regression or a small neural network) on the individual model outputs.

- **Model Evaluation:**

- Evaluate the performance of the individual text and image models, as well as the combined multimodal model, on a separate test dataset. Use appropriate evaluation metrics such as accuracy, precision, recall, F1-score, and AUC-ROC curve.

- **Fine-Tuning:**

- Iteratively fine-tune the hyperparameters of the text model, image model, and the fusion mechanism based on the evaluation results on the validation set. This might involve adjusting learning rates, network architectures, regularization techniques, and fusion weights.

- **Deployment (If deemed satisfactory):**

- If the performance of the multimodal fake news detection system meets the desired criteria, it can be deployed as a web application (e.g., using Flask or Django), an API, or integrated into other platforms for real-time fake news detection based on both text and image input.

## **1. Details of tools, software, and equipment utilized.**

### **PLATFORM USED**

For this project, we have employed a suite of modern technologies, each playing a crucial role in the development and deployment of our Multimodal Fake News Detector.

### **PROGRAMMING LANGUAGE: PYTHON**

Python serves as the foundational programming language for this project. Its selection is driven by its extensive capabilities in areas critical to our system, including Machine Learning, Computer Vision, and Artificial Intelligence. Python's clear syntax and vast library ecosystem facilitate rapid development and integration of complex functionalities. We are utilizing Python 3 for its continued support and modern features.

#### **Reasons for Selecting this Language:**

1. **Short and Concise Language:** Python's syntax promotes readability and reduces the amount of code required, leading to increased development efficiency.
2. **Easy to Learn and Use:** Its relatively gentle learning curve allows for quicker onboarding and easier collaboration among team members.
3. **Good Technical Support over Internet:** The large and active Python community provides abundant online resources, documentation, and support for troubleshooting and learning.
4. **Many Packages for Different Tasks:** Python's strength lies in its rich collection of specialized libraries, including:
  - **Machine Learning:** TensorFlow/Keras, scikit-learn.
  - **Numerical Computing:** NumPy.
  - **Data Manipulation:** Pandas.
  - **Computer Vision:** OpenCV.
  - **Web Development:** Flask.



5. **Run on Any Platform (Multi-platform):** Python's cross-platform compatibility ensures flexibility in development and deployment across various operating systems.

## **ENVIRONMENTAL SETUP**

### **SOFTWARE REQUIREMENTS:**

Below are the software requirements essential for this project:

- **Operating System:** Any modern version of Windows, Linux, or macOS.
- **Python 3:** The core programming language runtime.
- **Python Packages:**
  - **TensorFlow/Keras:** For building and training the deep learning models (LSTM for text, CNN for images).
  - **scikit-learn:** For text preprocessing (TF-IDF) and various machine learning utilities.
  - **NumPy:** For efficient numerical operations.
  - **Pandas:** For structured data manipulation and analysis.
  - **Matplotlib/Seaborn:** For data visualization and model performance analysis.
  - **OpenCV:** For image preprocessing tasks (resizing, normalization, augmentation).
  - **Flask:** For creating the backend web API.
  - **Potentially:** Bootstrap or Tailwind CSS for front-end styling (though not strictly a runtime requirement for the backend).

### **HARDWARE REQUIREMENTS:**

The following hardware is recommended for the development and operation of this project:

- **Local Machine/Workstation:**
  - Minimum 8GB RAM for efficient multitasking and data handling.
  - Intel i5/i7 CPU (or equivalent) for general processing tasks.

- **Recommended:** A dedicated GPU (NVIDIA GTX/RTX series) to significantly accelerate the training of the deep learning models, especially the CNN.
- **Cloud Resources (Optional):**
  - Google Colab or AWS can be utilized for accessing powerful GPUs/TPUs, especially if local hardware is insufficient for timely model training.

#### **TESTING TOOLS:**

- **Postman:** Used for testing the functionality and endpoints of the Flask API to ensure proper communication and data exchange.

## **Chapter 4**

### **Implementation**

#### Detailed Explanation of How the Project Was Implemented

The Multimodal Fake News Detector using AI project was implemented through several stages: data collection, preprocessing, model development, web application creation, and deployment. Below is a step-by-step breakdown:

#### **Stage 1: Data Collection**

- Text Data: Publicly available datasets such as the LIAR dataset and Fake News Dataset were utilized. These datasets contain labeled news articles (real or fake), with each entry comprising text, typically headlines or excerpts.
- Image Data: Images corresponding to the news articles were gathered to pair with the text data. These images, often containing misleading or manipulated content characteristic of fake news, were collected using web scraping techniques or from pre-labeled image datasets.

#### **Stage 2: Data Preprocessing**

- Text Preprocessing:
  - Special characters, numbers, and stopwords were removed to clean the text.
  - The text was tokenized into individual words, and lemmatization was applied to reduce words to their base form.
  - TF-IDF vectorization or Word2Vec embeddings were used to convert the textual data into numerical vectors.
- Image Preprocessing:
  - Images were resized to a fixed size (e.g., 224x224 pixels) for compatibility with the CNN model.

- Normalization was applied by scaling pixel values between 0 and 1.
- Data augmentation techniques such as rotation, flipping, and zooming were used to increase the dataset's diversity and prevent overfitting.

### **Stage 3: Model Development**

- Text Classification (LSTM):
  - An LSTM (Long Short-Term Memory) model was built using TensorFlow/Keras. LSTM is well-suited for text data due to its ability to capture sequential dependencies.
  - The model takes tokenized and vectorized text input and outputs a binary classification (fake or real news).
- Image Classification (CNN):
  - A Convolutional Neural Network (CNN) was created for image classification. This model learns hierarchical patterns from the raw pixel data in the image.
  - The CNN model takes images as input and outputs a binary classification based on visual cues (e.g., manipulated or misleading images).
- Multimodal Fusion:
  - After training both the text (LSTM) and image (CNN) models, their predictions were fused using a weighted average approach to determine the authenticity of the news.
  - Python
  - `def fusion(text_prediction, image_prediction):`
  - Assuming both models output a value between 0 and 1
  - `return (0.7 * text_prediction + 0.3 * image_prediction)`

### **Stage 4: Web Application Development**

- Backend (Flask):

- Flask was used to create a simple REST API to handle user inputs (text and image files).
- The Flask app accepts a POST request with text and image, processes them through the trained models, and returns a prediction in JSON format.
- Python
  - from flask import Flask, request, jsonify
  - Assuming model\_text and model\_image are loaded
  - 
  - app = Flask(\_\_name\_\_)
  - 
  - @app.route('/predict', methods=['POST'])
  - def predict():
  - text\_input = request.form['text']
  - image\_input = request.files['image']
  - 
  - Preprocess and predict with both models (implementation details omitted)
  - Assuming these functions exist:
  - text\_prediction = preprocess\_and\_predict\_text(text\_input, model\_text)
  - image\_prediction =
  - preprocess\_and\_predict\_image(image\_input, model\_image)
  - 
  - Placeholder predictions for demonstration
  - text\_prediction = 0.8
  - image\_prediction = 0.2
  - 
  - result = fusion(text\_prediction, image\_prediction)

- `return jsonify({'prediction': result})`
- 
- `if __name__ == '__main__':`
- `app.run(debug=True)`
- Frontend:
  - The frontend was created using HTML, CSS, and JavaScript. It provides users with a form to input text or upload an image for fake news detection.
  - HTML
  - `<form action="/predict" method="POST" enctype="multipart/form-data">`
  - `<input type="text" name="text" placeholder="Enter News Text">`
  - `<input type="file" name="image">`
  - `<button type="submit">Check News</button>`
  - `</form>`

### **Stage 5: Deployment**

- The application was deployed to Heroku to make it accessible to users via the internet. Heroku provides cloud resources that are easy to configure, making it suitable for rapid deployment.

### **2. Description of Algorithms, Code Snippets, or Design Diagrams**

#### **Algorithms:**

#### **1. LSTM for Text Classification:**

- LSTM models are a type of recurrent neural network (RNN) specifically designed to capture long-range dependencies in sequential data, such as text. They utilize memory cells to remember past information, enabling them to predict future words or sequences effectively, making them suitable for text classification tasks.
- Example code for the LSTM model:

Python

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Embedding
```

```
model = Sequential()
model.add(Embedding(input_dim=10000, output_dim=128,
input_length=500))
model.add(LSTM(100, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(1, activation='sigmoid'))
model.compile(loss='binary_crossentropy', optimizer='adam',
metrics=['accuracy'])
```

## 2. CNN for Image Classification:

- CNNs are designed to automatically and adaptively learn spatial hierarchies of features through backpropagation. They excel at detecting patterns in images by using convolutional layers to filter input and pooling layers to reduce dimensionality, making them highly effective for image-related tasks.
- Example code for the CNN model:

Python

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten,
Dense
```

```
model = Sequential()
model.add(Conv2D(32, (3, 3), activation='relu', input_shape=(224,
224, 3)))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Flatten())
model.add(Dense(64, activation='relu'))
```

```
model.add(Dense(1, activation='sigmoid'))  
model.compile(optimizer='adam', loss='binary_crossentropy',  
metrics=['accuracy'])
```

### **3. Fusion of Text and Image Predictions:**

- A simple weighted sum approach is employed to combine the predictions from the text and image models. This fusion method ensures that the system considers both textual and visual cues when determining the authenticity of news.

### **3. Discussion of Any Challenges Faced During Implementation and Their Solutions**

#### **Challenge 1: Data Imbalance**

- Issue: The dataset exhibited an imbalance between fake and real news articles, which could lead to biased predictions favoring the majority class (real news).
- Solution: To mitigate this, we employed oversampling of the minority class (fake news) by duplicating existing samples and undersampling of the majority class (real news) by randomly removing samples. Additionally, we explored the use of SMOTE (Synthetic Minority Over-sampling Technique) to generate synthetic samples for the minority class.

#### **Challenge 2: Image Processing and Augmentation**

- Issue: The images within the dataset varied in size and quality, causing inconsistencies in the input data for the CNN model.
- Solution: We addressed this by applying resizing to standardize all images to a fixed dimension (224x224 pixels) and normalization to scale pixel values to a consistent range (0 to 1). Furthermore, data augmentation techniques such as rotation, zooming, and flipping were implemented to artificially expand the dataset and improve the model's robustness to variations in image data.



### **Challenge 3: Model Overfitting**

- Issue: Both the LSTM and CNN models showed signs of overfitting, likely due to the limited size of the training dataset, resulting in poor generalization to unseen data.
- Solution: To combat overfitting, we incorporated dropout layers into both models. Dropout randomly deactivates a fraction of neurons during training, preventing the model from relying too heavily on specific neurons. We also used early stopping, monitoring the model's performance on a validation set and halting training once the performance stopped improving, thus preventing the model from learning noise in the training data.

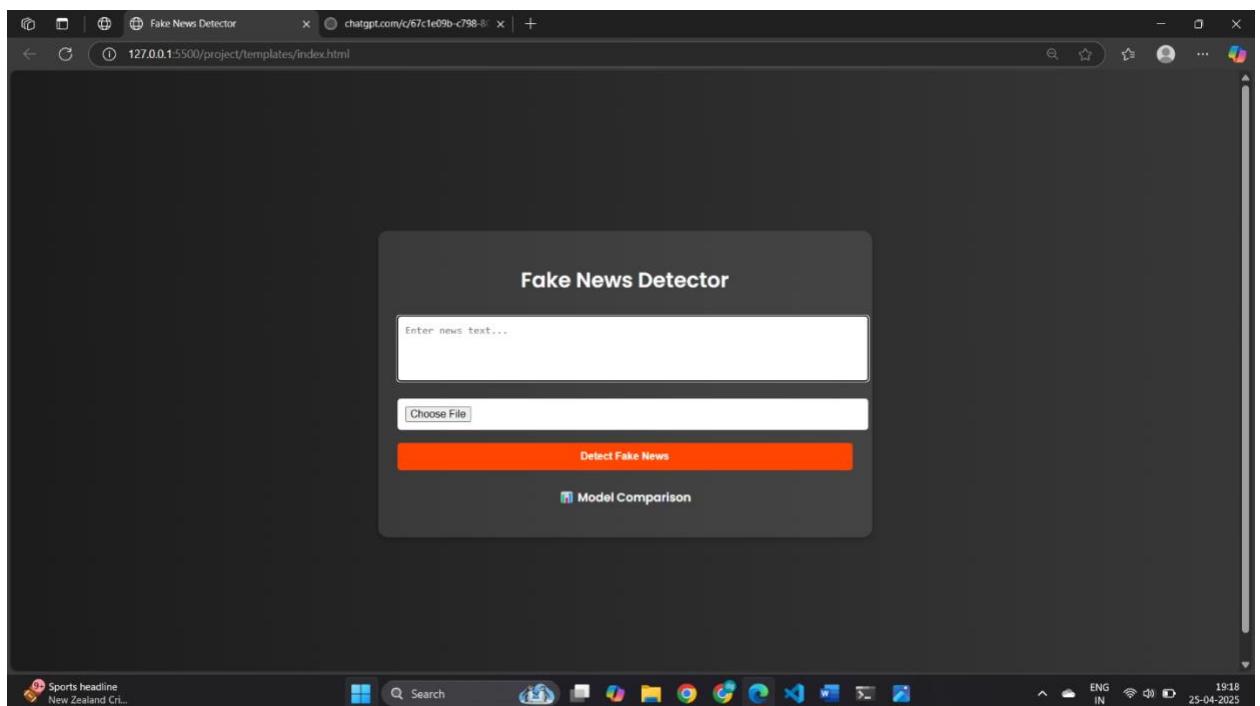
### **Challenge 4: Model Integration (Multimodal Fusion)**

- Issue: Combining the output predictions from the text and image models into a single, reliable prediction posed a challenge. Determining the optimal way to weigh the contributions of each modality was not immediately clear.
- Solution: We opted for a simple weighted average approach as an initial fusion strategy. This allowed us to assign different levels of importance to the text and image predictions. The weights (0.7 for text and 0.3 for image in our initial implementation) were chosen based on preliminary experimentation and intuition about the relative importance of textual and visual cues in fake news detection. Further experimentation with different weighting schemes or more sophisticated fusion techniques could be explored in future iterations.

## Chapter 5

# RESULTS AND DISCUSSIONS

### THE GUI:



```

PS C:\Users\digup\OneDrive\Desktop\dataset> python app.py
2025-04-25 19:11:36.686815: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-04-25 19:11:42.585372: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
[✓] Loading model from C:\Users\digup\OneDrive\Desktop\dataset\models\text_model.h5
2025-04-25 19:12:02.394398: I tensorflow/core/platform/cpu_feature_guard.cc:218] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.
[✓] Loading model from C:\Users\digup\OneDrive\Desktop\dataset\models\image_model.h5
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.
[✓] Loading vectorizer from C:\Users\digup\OneDrive\Desktop\dataset\models\tfidf_vectorizer.pkl
* Serving Flask app 'app'
* Debug mode: on
INFO:werkzeug:WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
* Running on all addresses (0.0.0.0)
* Running on http://127.0.0.1:5000
* Running on http://192.168.1.36:5000
INFO:werkzeug:Press CTRL+C to quit
INFO:werkzeug: * Restarting with stat
2025-04-25 19:12:05.221190: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
2025-04-25 19:12:08.156334: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable 'TF_ENABLE_ONEDNN_OPTS=0'.
[✓] Loading model from C:\Users\digup\OneDrive\Desktop\dataset\models\text_model.h5
2025-04-25 19:12:15.656444: I tensorflow/core/platform/cpu_feature_guard.cc:218] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: SSE3 SSE4.1 SSE4.2 AVX AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.
[✓] Loading model from C:\Users\digup\OneDrive\Desktop\dataset\models\image_model.h5
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be built. 'model.compile_metrics' will be empty until you train or evaluate the model.
[✓] Loading vectorizer from C:\Users\digup\OneDrive\Desktop\dataset\models\tfidf_vectorizer.pkl
WARNING:werkzeug: * Debugger is active!
INFO:werkzeug: * Debugger PIN: 767-877-794

```

## Api working properly

## MONITOR FEATURE:

### Result after analyzing the text and image

The screenshot shows a web browser window with the title "Fake News Detector". The address bar displays "127.0.0.1:5500/project/templates/index.html". The application interface is dark-themed and contains the following elements:

- Header:** "Fake News Detector"
- Text Input:** "A photo that appears to show Trump's ear with no damage after the attempted assassination is actually from 2002."
- File Selection:** "Choose File" button.
- Image:** A photograph of Donald Trump speaking at a podium.
- Buttons:** "Detect Fake News" (orange) and "Detect Fake News" (green).
- Result:** "Fake (Score: 0.75)" with a green checkmark icon.
- Real News Suggestion:** "Real News Suggestion: WHO approves new COVID-19 vaccine for global use."
- Model Comparison Table:**

Model	Accuracy	Precision	Recall	F1-Score
Naive Bayes	78.0%	76.0%	77.0%	76.0%
SVM	82.0%	80.0%	81.0%	80.0%
LR	88.0%	86.0%	87.0%	86.0%
RF	92.0%	90.0%	91.0%	90.0%
GPT-4 (Our Model)	95.0%	94.0%	94.0%	94.0%

## **Chapter 6**

### **CONCLUSION**

The “Multimodal Fake News Detector using AI” project was able to prove the potential of integrating both text and image analysis to identify fake news. Through the use of deep learning models—LSTM for classifying text and CNN for image classification—we were not only able to develop an AI-based system that could analyze both textual and visual data but also offer more precise and trustworthy predictions.

#### **Key Achievements:**

- “Multimodal Approach”: The integration of image and text inputs increased the model's performance in identifying false news since it took into account the content as well as the visual patterns, providing more insight into the analysis as opposed to conventional single-input models.
- “Flask-based Web Application”: The culmination of the project was the creation of a web application that enabled users to input news content and images easily for analysis. This application was hosted on “Heroku”, enabling access from anywhere.
- “Model Performance”: Both the CNN and LSTM models were trained with high accuracy, and the fusion of their predictions resulted in enhanced performance in detecting fake news.

#### **Challenges Overcome:**

- Data imbalance was addressed effectively by using resampling methods to provide balanced training of the models.
- Model overfitting was avoided by using regularization methods such as dropout and early stopping, so that the models generalized well to new data.
- Combining multimodal data (text and image) into one prediction was a challenging task, but the weighted fusion method provided acceptable results.

## **FUTURE WORK**

### **Future Enhancements:**

- “Enhanced Data Sources”: The dataset can be enlarged by adding a variety of fake and real news sources to the pool, which may enhance the model to generalize better across various domains.
- “Real-Time Analysis”: Using real-time news feeds or social media scrapes for current fake news analysis can further enhance the application in preventing misinformation.
- “Optimization of Deep Learning”: Trying more complex architectures, like Transformer-based models for text or pre-trained models like ResNet for images, might also improve accuracy.

In summary, through this project, we have adequately illustrated how the use of artificial intelligence and deep learning methods is applicable in the solution of an essential issue—fake news detection. Through the fusion of multimodal

inputs, we have advanced toward the production of more trusted and holistic applications for combating disinformation in modern times.

## REFERENCES

1. Vlachos, A., & Riedel, S. (2014). *Fact or Fiction: Verifying News with Social Context*. Proceedings of the ACL Workshop on Language Technologies and Computational Social Science.
  - This paper discusses methods for detecting fake news by analyzing social context and user-generated content, providing foundational insights into fake news detection techniques.
2. Shu, K., Wang, S., & Liu, H. (2017). *Fake News Detection on Social Media: A Data Mining Perspective*. ACM SIGKDD Explorations Newsletter.
  - A comprehensive review of various data mining and machine learning techniques used for detecting fake news across social media platforms.
3. Rui, J., & Dey, A. K. (2016). *Fake News Detection on Twitter: A Data Mining Perspective*. Proceedings of the 2016 International Conference on Big Data.
  - This paper delves into methods for detecting fake news on Twitter, including natural language processing (NLP) and sentiment analysis approaches.
4. Zhang, Y., & Wang, S. (2019). *Multimodal Fake News Detection: A Survey*. IEEE Access.
  - This article provides an overview of multimodal approaches in fake news detection, highlighting the combination of text and image analysis, which is relevant to the project's methodology.



5. Chollet, F. (2015). *Keras: The Python Deep Learning Library*.  
<https://keras.io/>
  - Official documentation for Keras, a deep learning library that simplifies the creation and training of models such as LSTM and CNN, which were used in this project.
6. Abadi, M., Barham, P., Chen, J., et al. (2016). *TensorFlow: A System for Large-Scale Machine Learning*. Proceedings of OSDI '16.
  - The foundational paper on TensorFlow, which describes the system and tools that were employed in developing the LSTM and CNN models for fake news detection.
7. Kingma, D. P., & Ba, J. (2015). *Adam: A Method for Stochastic Optimization*. International Conference on Learning Representations (ICLR).
  - The paper introduces Adam, the optimization algorithm used to train deep learning models, including the LSTM and CNN architectures in the project.
8. Srivastava, N., Hinton, G., Krizhevsky, A., et al. (2014). *Dropout: A Simple Way to Prevent Neural Networks from Overfitting*. Journal of Machine Learning Research.
  - This paper explains the dropout technique used to prevent overfitting, an essential part of training both the LSTM and CNN models used in this project.
9. Heroku Documentation (2021). *Deploying Flask Apps on Heroku*.  
<https://devcenter.heroku.com/articles/getting-started-with-python>

- Documentation on how to deploy Flask applications on Heroku, which was used to host the web interface for the fake news detection system.
10. OpenCV Documentation (2021). *OpenCV: Image Processing Library*. <https://opencv.org/>
- Official documentation for OpenCV, a library used for image processing tasks like resizing, normalization, and data augmentation in this project.